

1 File systems

A *user view* that defines a file system in terms of the abstractions that the operating system provides. An **implementation view** that defines the file system in terms of its **low level implementation**

1.1 User view

Important aspects of the user view include:

- The file **abstraction** which hides away **implementation details** to the user (similar to processes and memory)
- **File naming** policies (abstracts storage details), user file attributes (e.g. size, protection, owner, protection, dates)
- There are also **system attributes** for files (e.g. non-human readable file descriptors (similar to a PID), archive flag, temporary flag, etc.)
- Directory structures and organisation
- System calls to interact with the file system

1.2 Structures

1.2.1 Overview

Different directory structures have been used over the years

- Single level: all files in the same directory (reborn in consumer electronics)
- Two or multiple level directories (hierarchical): tree structures
- Absolute path name: from the root of the file system Relative path name: the current working directory is used as the starting point
- Directed acyclic graph (DAG): allows files to be shared (i.e. links to files or sub-directories) but cycles are forbidden
- Generic graph structure in which links and **cycles can exist**

The use of **DAGs** and **generic graph** structures results in significant complications in the implementation

1.2.2 DAG and graph complications

When searching the file system:

- Cycles can result in **infinite loops**.
- Sub-trees can be traversed **multiple times**.

Files have **multiple** absolute file names. Deleting files becomes a lot more **complicated** (i.e. links may no longer point to a file, **inaccessible cycles** may exist). A *garbage collection* scheme may be required to remove files that are **no longer accessible** from the file system tree (that are part of a cycle only)

1.3 Directories

Directories contain a list of **human readable** file names that are mapped onto unique identifiers and disk locations. They provide a **mapping** of the **logical** file onto the **physical** location. Retrieving a file comes down to searching a directory file as fast as possible:

- A simple random order of directory entries might be **insufficient** (search time is linear as a function of the number of entries)
- Indexes or *hash tables* can be used
- They can store all file related attributes (e.g. file name, disk address Windows) or they can contain a **pointer** to the data structure that **contains** the details of the file (Unix)

1.3.1 System calls

Similar to files, directories are manipulated using system calls:

- create/delete: a new directory is created/deleted
- opendir, closedir: add/free directory to/from internal tables
- readdir, return the next entry in the directory file
- Others: rename, link, unlink, list, update

Directories are special **files** that **group files** together and of which the structure is defined by the file system. **A bit is set to indicate that they are directories!**

2 Files

2.1 Types

Many OSs support **several** types of file. Both Windows and Unix (including OS X) have regular files and directories:

- Regular files contain user data in ASCII or binary (well defined) format
- Directories group files together (but are files on an implementation level)

Unix also has character and block special files:

- **Character** special files are used to model **serial I/O devices** (e.g. keyboards, printers)
- **Block** special files are used to model, e.g. **hard drives**

3 System calls

3.1 Types

File control blocks (**FCBs**) are kernel data structures, i.e. they are **protected** and only accessible in *kernel mode*!

Allowing user applications to access them directly could **compromise** their integrity. System calls enable a user application to **ask** the operating system to carry out an action on its behalf (in kernel mode).

There are two different categories of system calls:

- File manipulation: `open()`, `close()`, `read()`, `write()`, ...
- Directory manipulation: `create()`, `delete()`, `readdir()`, `rename()`, `link()`, `unlink()`, `list()`, `update()`

4 Implementation context

Irrespective of the type of file system, a number of **additional** considerations have to be addressed, including:

- Disk partitions, partition tables, boot sectors, etc.
- Free space management (free memory)
- System wide and per process file tables (process tables)

Low level formatting **writes** sectors to the disk, high level formatting **imposes a file system** on top of this (using blocks that can cover multiple sectors).

5 Hard Disk Structures

Disks are usually **divided** into multiple partitions. An **independent** file system may exist on each partition. **Master boot record** is located at start of the entire drive:

- Used to **boot** the computer (BIOS reads and executes MBR)
- Contains *partition table* at its end with active partition
- One partition is listed as **active** containing a boot block to **load the operating system**



Figure: Layout of a Disk

6 Partition layouts

The layout of a partition differs depending on the file system. A UNIX partition contains:

- The partition boot block: Contains **code to boot** the operating system. **Every** partition has boot block **even if** it does not contain OS
- Super block contains the **partitions details**, e.g., partition size, number of blocks, I-node table size
- Free space management contains, e.g., a bitmap or linked list that indicates the free blocks
- I-nodes: an array of data structures, one per file, telling all about the files
- Root directory: the top of the file-system tree
- Data: files and directories



Figure: Layout of a Partition

7 Disk space management

Two methods are commonly used to keep track of free disk space: *bitmaps* and *linked lists*. Note that these approaches are very similar to the ones to keep track of **free memory**.

- Bitmaps represent each block by a **single bit** in a map
- The size of the bitmap **grows with the size** of the disk but is **constant** for a given disk
- Bitmaps take comparably **less space** than linked lists

A Linked List of disk blocks (also known as **Grouping**).

- We use **free blocks** to hold the **numbers of the free blocks** (hence, they are no longer free). E.g. with a 1KB block a 32-bit disk block number, each block will hold 255 free blocks (one for the pointer to the next block). Since the free list **shrinks** when the disk becomes full, this **is not wasted space**.
- Blocks are linked together, i.e., multiple blocks list the free blocks
- The size of the list **grows** with the size of the disk and **shrinks** with the size of the blocks
- Linked lists can be modified by keeping track of the number of **consecutive free blocks** for each entry (known as Counting)

7.1 Comparison

Bitmaps:

- Require extra space. E.g: If block size = 212 bytes and disk size = 230 bytes (1 GB) bitmap size: $230/2^{12} = 218$ (32KB)
- Keeping it in main memory is possible only for **small disks**.

Linked lists:

- **No waste** of disk space
- We only need to keep in memory **one block** of pointers (load a new block when need)

8 File tables

Apart from the free space memory tables, there is a number of key data structures stored in memory:

- An in-memory mount table
- An in-memory directory cache of **recently** accessed directory information
- A **system-wide** open file table, containing a copy of the FCB for every currently open file in the system, including location on disk, file size, and open count (number of processes that use the file)
- A **per-process** open file table, containing a pointer to the system open file table

Reference section

user view

The user view defines how the file system looks like to regular users (and programmers) and relates to abstractions