

# Contents

<b>1</b>	<b>Reference Monitor</b>	<b>2</b>
<b>2</b>	<b>Placement</b>	<b>2</b>
2.1	Lower Is Better . . . . .	2
<b>3</b>	<b>OS Integrity</b>	<b>2</b>
3.1	Modes of operation . . . . .	2
3.2	Controlled Invocation . . . . .	2
3.2.1	Controlled Invocation: Interrupts . . . . .	2
<b>4</b>	<b>Descriptors and Selectors</b>	<b>3</b>
4.1	Interrupt-gates . . . . .	3

# 1 Reference Monitor

An access control concept that refers to an abstract machine that **mediates all access** to objects by subjects

- Must be tamper proof / resistant
- Must **always be invoked** when access to an object is required
- Must be small enough to be verifiable / subject to analysis to ensure correctness

## 2 Placement

Can be placed anywhere within the system, a variety of locations relative to the program being run

- Hardware Dedicated registers for defining privileges
- Operating system kernel E.g. Virtual Machine Hypervisor
- Operating system Windows security reference monitor
- Services Layer JVM, .NET
- Application Layer Firewalls

### 2.1 Lower Is Better

Using a reference monitor or other security features at a **lower level** means:

- We can assure a higher degree of security
- Usually simple structures to implement
- Reduced performance overheads
- Fewer layer below attack possibilities

However: **access control decisions** are far removed from applications

## 3 OS Integrity

The operating system: arbitrates access requests, is itself an resource that must be accessed. This is a **conflict**, we want to **use** the OS but **not mess with it**.

### 3.1 Modes of operation

Defines which actions are permitted in which mode, e.g. system calls, machine instructions, I/O. Distinguish between computations done on behalf of: **the OS and the user**. A *status flag* within the CPU allows the OS to operate in different modes

### 3.2 Controlled Invocation

Many functions are held at kernel level, but are quite reasonably called from within user level code

- Network and File IO
- Memory allocation
- Halting the CPU (at shutdown only!)

We need a mechanism to transfer between kernel mode (ring 0 - root) and user mode (ring 3 - user)

#### 3.2.1 Controlled Invocation: Interrupts

Also known as exceptions / traps. In many ways is the **hardware equivalent** to a **software exception** not always bad. Handled by an *interrupt handler* which resolves the issue and returns to the original code. Given an interrupt, the CPU will **switch execution** to the location given in an *interrupt descriptor table*

## 4 Descriptors and Selectors

- Descriptors hold information on crucial system objects like **kernel structure locations**
- Descriptors are held in *descriptor tables*
- Contain a Descriptor Privilege Level (DPL)
- Descriptors are indexed by selectors
- Loaded when required, e.g. on jump calls
- The CPU protects the kernel by **checking** the Current Privilege Level (CPL) when a Selector is loaded

### 4.1 Interrupt-gates

- The code segment (CS) register in x86 CPUs has 2 bits reserved for the Current Privilege Level (CPL)
- Descriptors that have a privilege level **higher than where they point are called gates**
- Since these descriptors are **created by the kernel**, they offer a **secure means of entry into ring 0**

### 4.2 Patching the kernel

If you can run custom PL 0 code (compromised driver?), you can insert your own handler *Rootkit*

## 5 Processes and Threads

- A process is a program being executed. Important unit of control:
  - Exists in its own address space
  - Communicates with other processes via the OS
  - Separation for security
- A Thread is a strand of execution within a process. Shares a common address space

## 6 Memory Protection

Segmentation divides data into logical units

- Good for security
- Challenging memory management
- Not used much in modern OSs

Paging divides memory into pages of equal size

- Efficient memory management
- Worse for access control
- Extremely common in modern OSs

## 7 Side channel exploits

In most operating systems, the entire kernel is stored in the **upper address space**. Pages in this area are flagged as supervisor, and cannot be accessed in ring 0. In Intel (and some other) CPUs, its common to **speculatively evaluate** code prior to reaching it. If the branch isn't taken, changes are rolled back, however cache isn't.

### 7.1 Meltdown

Meltdown breaks the mechanism that keeps applications from accessing arbitrary system memory. Consequently, applications can access system memor

## 7.2 Spectre

Spectre tricks other applications into accessing arbitrary locations in their memory

## Reference section

placeholder